

# Microeconometrics Task - Problem Set 3

Ricardo Semião e Castro

2024-06-06

## 1 Setup

The packages used were:

```
library(nprobust)
library(gt)
library(furrr)
library(tidyverse)

set.seed(0306152529)
theme_set(theme_bw())
plan(multisession, workers = 7)
```

Lets define a function to plot the results from the models in this questions:

```
plot_lp <- function(x, tau_col = "tau.us", se_col = "se.rb") {
  ggplot(data, aes(x1, f)) +
    geom_line() +
    geom_line(aes(eval, .data[[tau_col]]), x, linetype = "dashed") +
    geom_ribbon(
      aes(
        eval, .data[[tau_col]],
        ymin = .data[[tau_col]] - .data[[se_col]],
        ymax = .data[[tau_col]] + .data[[se_col]]
      ),
      x, linetype = "dashed", fill = NA, color = "blue"
    )
}
```

## 2 Question 1 and 2

We can create the simulated dataset with the following function:

```
sim_data <- function(n, beta1, beta2) {
  e <- replicate(2, rnorm(n))

  x <- mvtnorm::rmvnorm(n, sigma = matrix(c(1, 0.9, 0.9, 1), 2, 2)) %>%
    apply(2, pnorm)

  f <- sin(beta1 * x[,1])

  y <- cbind(f + e[,1], f + beta2 * x[,2] + e[,2])
}
```

```
cbind(x, y, f) %>%
  `colnames<-`(c("x1", "x2", "y1", "y2", "f")) %>%
  .[order(.[, "x1"]), ]
}
```

We can now define the parameters and create the simulated dataset:


```
n <- 10000

beta1 <- 4
beta2 <- 2

data <- sim_data(n, beta1, beta2)
```

And finally, plot the function  $\sin(\beta_1 * x_1)$ :

```
ggplot(data, aes(x1, f)) +
  geom_line() +
  labs(title = "Plot of f", y = "f", x = "X1")
```



## 2.1 Question 3 and 4

We can use the `lproburst` package/function to estimate a local polynomial on  $Y_1$ , using the data from  $X_1$ .

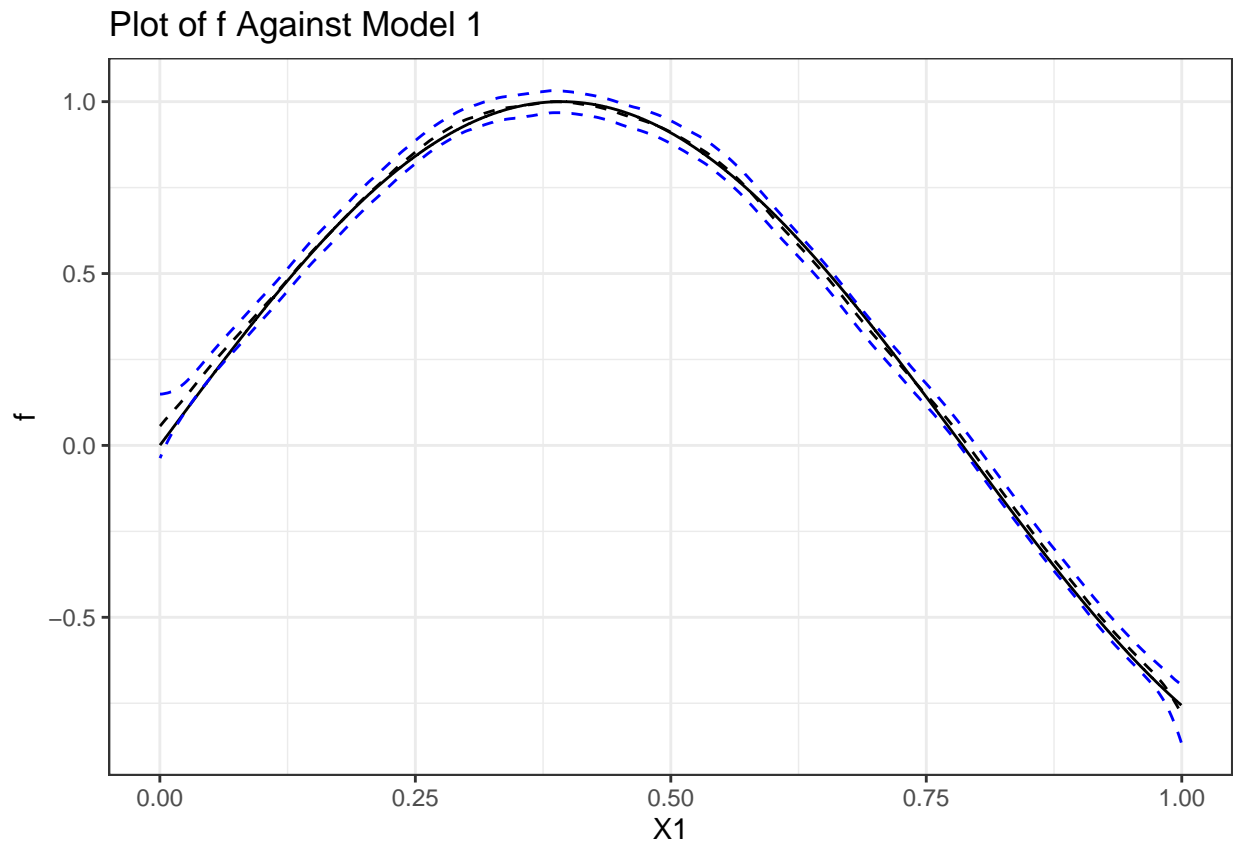
The specification of the model considers:

- `n` points of evaluation, related to the parameter `neval`.

- Polynomial order of 0, related to the parameter `p`. A level of zero is related to a “simple” local linear regression, as wanted.
- Derivative order of 0, related to the parameter `deriv`. To estimate a regression function, as wanted.
- Bandwidth  $h$  selected using `lpbwselect`, following the method of Calonico, Cattaneo and Farrell (2018) – “second-generation DPI implementation of IMSE-optimal bandwidth”. The same was used for the bias correcting bandwidth  $b$ .
- Kernel: Epanechnikov kernel.

```
mod1 <- lpobust(data[, "y1"], data[, "x1"], neval = n)

plot_lp(mod1$Estimate) +
  labs(title = "Plot of f Against Model 1", y = "f", x = "X1")
```



Besides the visualization, we can get statistics for the distance between the function and its estimated counterpart. There are several options, for example:

```
agg_funs <- list(
  `Mean abs.` = mean,
  `Mean sqr.` = \(x) mean(x^2),
  `Max abs.` = max,
  `Median abs.` = median
)
```

Consider the function below to extract these measures of difference from the model. Note that we might also be interested in considering:

- The difference between the true and the estimate.
- An lower bound, using the lowest difference among the whole confidence interval, at each  $x$ .

- An upper bound, using the highest difference among the whole confidence interval, at each  $x$ .

```
diff_lp <- function(x, tau_col = "tau.us", se_col = "se.rb") {
  n <- nrow(x)

  values <- cbind(
    x[, tau_col] - x[, se_col] - data[, "f"],
    x[, tau_col] - data[, "f"],
    x[, tau_col] + x[, se_col] - data[, "f"]
  )

  imap_dfr(agg_funs, function(fun, name) {
    list(
      Measure = name,
      Estimate = fun(abs(values[, 2])) / n,
      Lower = fun(apply(values, 1, \ (x) min(abs(x)))) / n,
      Upper = fun(apply(values, 1, \ (x) max(abs(x)))) / n
    )
  })
}
```

Lets apply it to our model 1:

```
diff_lp(mod1$Estimate) %>%
  gt() %>%
  fmt_scientific(-1, decimals = 2)
```

Measure	Estimate	Lower	Upper
Mean abs.	$1.19 \times 10^{-6}$	$9.18 \times 10^{-7}$	$4.65 \times 10^{-6}$
Mean sqr.	$2.25 \times 10^{-8}$	$1.09 \times 10^{-8}$	$2.34 \times 10^{-7}$
Max abs.	$5.60 \times 10^{-6}$	$3.77 \times 10^{-6}$	$1.49 \times 10^{-5}$
Median abs.	$1.03 \times 10^{-6}$	$9.32 \times 10^{-7}$	$4.37 \times 10^{-6}$


We can see that the values are very similar, which is expected, as indeed,  $Y_1$  is explained only by  $X_1$  and some random variation ( $\epsilon_1$ )

## 2.2 Question 5 and 6

The model specification is the same, but now we use  $Y_2$  against  $X_1$ .

```
mod2 <- lprobust(data[, "y2"], data[, "x1"], neval = n)

plot_lp(mod2$Estimate) +
  labs(title = "Estimation of Model 2", y = "f", x = "X1")
```



And the differences statistics:

```
diff_lp(mod2$Estimate) %>%
  gt() %>%
  fmt_scientific(-1, decimals = 2)
```

Measure	Estimate	Lower	Upper
Mean abs.	$9.90 \times 10^{-5}$	$9.56 \times 10^{-5}$	$1.03 \times 10^{-4}$
Mean sqr.	$1.26 \times 10^{-4}$	$1.19 \times 10^{-4}$	$1.33 \times 10^{-4}$
Max abs.	$2.03 \times 10^{-4}$	$1.94 \times 10^{-4}$	$2.13 \times 10^{-4}$
Median abs.	$9.91 \times 10^{-5}$	$9.57 \times 10^{-5}$	$1.02 \times 10^{-4}$

We can see that the values are very different, which is expected, as indeed  $Y_2$  is not explained only by  $X_1$  through  $f$ , the estimated function is also capturing the effect of  $\beta_2 X_2$  in  $Y_2$  through its correlation with  $X_1$ . The estimation is indeed similar to a combination of a sin curve and a linear trend.

## 2.3 Question 7, 8, and 9

Following what was stated above, we need to cleanse the relation that  $X_1$  has on  $Y_2$  through  $\beta_2 X_2$ . For that, we will:

- Find the (non-linear) relation between  $X_2$  and  $X_1$ , and between  $Y_2$  and  $X_1$ .
- Then with the unexplained parts of such relations (the residuals), we can find the relation  $Y_2 \sim X_2$ .
- We can use that to “cleanse”  $Y_2$  from the effect from  $X_2$ .
- Lastly, we can use that adjusted  $Y_2$  to find a better estimation for  $f$ .

```

mod3_x2 <- lprobust(data[, "x2"], data[, "x1"], neval = n)
x2_res <- data[, "x2"] - mod3_x2$Estimate[, "tau.us"]


mod3_y2 <- lprobust(data[, "y2"], data[, "x1"], neval = n)
y2_res <- data[, "y2"] - mod3_y2$Estimate[, "tau.us"]

mod3_res <- lm(y2_res ~ 0 + x2_res)

y2_adj <- data[, "y2"] - data[, "x2"] * mod3_res$coefficients[1]
mod3 <- lprobust(y2_adj, data[, "x1"], neval = n)

plot_lp(mod3$Estimate) +
  labs(title = "Estimation of Model 3", y = "f", x = "X1")

```



And the differences statistics:

```

diff_lp(mod3$Estimate) %>%
  gt() %>%
  fmt_scientific(-1, decimals = 2)

```

Measure	Estimate	Lower	Upper
Mean abs.	$3.82 \times 10^{-6}$	$1.06 \times 10^{-6}$	$7.21 \times 10^{-6}$
Mean sqr.	$1.61 \times 10^{-7}$	$1.85 \times 10^{-8}$	$5.34 \times 10^{-7}$
Max abs.	$7.09 \times 10^{-6}$	$3.79 \times 10^{-6}$	$1.18 \times 10^{-5}$

---

Median abs.    $3.83 \times 10^{-6}$     $8.78 \times 10^{-7}$     $7.17 \times 10^{-6}$

We can see that the values are very similar. This is a sign that our logic of cleansing the effect of  $X_2$  was actually right, and we got a estimation only on the effect of  $X_1$  trough  $f$ .

The estimated  $\beta_2$  can be seen below. It is decently close to 2, as it should, following the Frisch-Waugh-Lovell-like result that was used in its estimation.

```
summary(mod3_res)$coefficients["x2_res", ]

##      Estimate      Std. Error      t value      Pr(>|t|)
## 2.051023e+00  7.580837e-02  2.705537e+01 1.210696e-155
```

## 2.4 Extra: Monte Carlo

As an extra exercise, I've created a possible Monte-Carlo solution. I'll use the `furrr` package for parallel computing. We just need to estimate each of the three models in a loop, saving the results.

As this was too computationally expensive, I did not run the full loop to get the results. Still, I did check if the code would be able to run correctly.

```
n_iter <- 1000

results <- future_map(seq_len(n_iter), function(iter) {
  data <- sim_data(n, beta1, beta2)

  mod1 <- lprobust(data[, "y1"], data[, "x1"], neval = n)

  mod2 <- lprobust(data[, "y2"], data[, "x1"], neval = n)

  mod3_param <- lprobust(data[, "x2"], data[, "x1"], neval = n)
  x2_res <- data[, "x2"] - mod3_param$Estimate[, "tau.us"]
  mod3_nparam <- lprobust(data[, "y2"], data[, "x1"], neval = n)
  y2_res <- data[, "y2"] - mod3_nparam$Estimate[, "tau.us"]
  mod3_res <- lm(y2_res ~ -1 + x2_res, neval = n)

  list(
    mod1 = mod1$Estimate,
    mod2 = mod2$Estimate,
    mod3 = mod3$Estimate,
    beta2 = mod3_res$coefficients["x2_res"]
  )
},
.options = furrr_options(seed = TRUE)
) %>%
transpose()
```

Now, we can calculate a `diff_lp` over all the models in one big matrix of results (`do.call(rbind, results$modX)`):

```
diff_lp(do.call(rbind, results$mod1)) %>%
  gt() %>%
  fmt_scientific(-1, decimals = 2)

diff_lp(do.call(rbind, results$mod2)) %>%
  gt() %>%
  fmt_scientific(-1, decimals = 2)
```

```
diff_lp(do.call(rbind, results$mod3)) %>%  
  gt() %>%  
  fmt_scientific(-1, decimals = 2)
```

For the beta, we can see, in average, how much it strays away from  $\beta_2$ .

```
mean(abs(list_c(results$beta2) - beta2))
```